

STATEMENT OF WORK · AUGUST 2026 · TIN DEVELOPMENT · SUPPLIER-PROTECTED EDITION

Canada Lighting Supplies — Spec Coverage & Catalog Operations Full-Build Ledger

Every commit, every build day, every dollar — reconstructed from git history and the committed plan/spec docs, priced three ways: our actual cost, the client rate (5×), and a traditional dev shop.

18

Commits priced

2

Coding days

0

Planning days

Aug 2026 – Aug 2026

Engagement span

HEADLINE

Three ways to read the cost

ACTUAL COST (PARTNER PRICE)

\$3,528

build + integration · our cost

BILLABLE (5x · CLIENT PRICE)

\$17,640

5x on coding · planning flat · incl. integration

TRADITIONAL DEV SHOP

\$35,842

266 human hours @ \$135/hr

Tunable plan: man-hours/day = commits × 1.6 × tier-weight, capped 1–38h (avg 18.6h) · human team ≤5h/day (flat, rates below) · billable = 5x actual on coding only · traditional = 5x our man-hours @ \$135/hr.

Labour cost by category

Core AI programming is separated from the human project team so each line is costed on its own basis — coding at tiered model rates, the team (senior + junior programmer, database expert, design expert, QA) flat at human rates.

LABOUR CATEGORY	HOURS	ACTUAL	BILLABLE	TRADITIONAL
Core AI programming 2 coding days · tiered model rates	37h	\$2,728	\$13,640	\$25,042
Human project team 0 non-coding days · 10 roles, flat	0h	\$0	\$0	\$0
Third-party integration 1 providers · API/SDK wiring	16h	\$800	\$4,000	\$10,800
Total labour	53h	\$3,528	\$17,640	\$35,842

LEGEND

Model tiers

Super Premium

TOP MODELS · HIGH EFFORT

Claude Opus 4.8 · GPT-5 (high reasoning)

ACTUAL

BILLABLE 5×

\$80/hr

\$400/hr

Cryptographic signing, novel architecture, multi-agent, refactors.

Premium

FLAGSHIP · STANDARD EFFORT

Claude Sonnet 4.5 · Gemini 2.5 Pro

ACTUAL

BILLABLE 5×

\$50/hr

\$250/hr

Standard features, APIs, integrations, bug fixes.

Standard

FAST · LOW EFFORT

Claude Haiku 4.5 · GPT-5 mini

ACTUAL

BILLABLE 5×

\$25/hr

\$125/hr

Templated screens, CRUD, config, scaffolding, docs.

Human project team (non-coding days, flat rate)

ROLE	RATE	DAYS	HOURS	COST	RESPONSIBILITY
Senior Programmer	\$150/hr	0	0h	\$0	Architecture, code review, hard problems, sequencing
Database Expert	\$140/hr	0	0h	\$0	Schema, migrations, RLS, query performance
Design Expert	\$135/hr	0	0h	\$0	UX flows, screen design, design system
QA Engineer	\$95/hr	0	0h	\$0	Test planning, TDD support, regression, sign-off
Junior Programmer	\$75/hr	0	0h	\$0	Templated screens, CRUD, wiring, fixes
Project Manager	\$110/hr	0	0h	\$0	Sprint planning, client comms, timelines, risk
Business Analyst	\$105/hr	0	0h	\$0	Requirements, acceptance criteria, user stories
Team total		0	0h	\$0	

ROLE	RATE	DAYS	HOURS	COST	RESPONSIBILITY
DevOps Engineer	\$145/hr	0	0h	\$0	CI/CD, cloud infra, monitoring, deployments
Security Engineer	\$160/hr	0	0h	\$0	Audits, pen testing, compliance, RLS, Stripe security
UX Researcher	\$120/hr	0	0h	\$0	Usability testing, user interviews, flows
Team total		0	0h	\$0	

ENGAGEMENT SCOPE

Phases from commit history

PHASE	VOLUME
General Platform Work	15 commits

EVIDENCE

Planning & specification artifacts

The build is backed by 4 committed specification documents (3,349 words total). These plans — not just the commits — define scope, sequencing, and acceptance criteria, and are the basis for the effort estimate below.

DOCUMENT	DEPTH
README.md	1,374 words
docs/source-material-system.md	1,223 words
docs/droid-pipeline-runbook.md	594 words
AGENTS.md	158 words

INTEGRATIONS

Third-party API & SDK integration

Third-party providers are shown by category only; names are withheld in this supplier-protected edition. Integration effort is derived from provider-tagged commits (auth, webhooks, checkout, tokens, error handling and provider testing), floored per provider, and priced at the Premium rate.

PROVIDER	COMMITTS	INTEG. HRS	ACTUAL	BILLABLE 5×	TRADITIONAL
Voice & SMS Provider	1	16h	\$800	\$4,000	\$10,800
Integration subtotal	1	16h	\$800	\$4,000	\$10,800

Edition: **Supplier-protected**. Integration totals are additive to the build ledger above (API/SDK wiring layered on feature commits).

How man-hours and the traditional multiple are derived

There is no honest single "AI makes you Nx faster" number. The evidence is mixed and measurement diverges sharply from perception, so this ledger does *not* assume a fast single developer. It models compression from **parallelism and verification discipline**, then benchmarks the same delivered scope against a traditional team.

FINDING	SOURCE	IMPLICATION HERE
Experienced devs on familiar codebases were 19% slower with AI, yet believed they were ~20% faster. The ~39-point perception gap is the single most important calibration signal: self-report inflates apparent AI gains by the same margin it inflates traditional-team estimates.	METR RCT, 16 devs / 246 real tasks	Common industry intuition places a traditional team at ~4x AI speed. Applying the METR perception correction ($\div 0.81$ for actual vs perceived) raises that to ~5x — the multiplier used here.
Copilot users finished a task in 21 min vs 2h41m (55% faster) on greenfield/boilerplate.	GitHub controlled study	Gains are real but concentrated in low-complexity work — our Standard tier.
Measured savings cluster at ~3.6 hrs/week ; self-reported savings run far higher (JetBrains 2025).	Industry surveys	Confirms modest, task-dependent single-threaded gains.
Juniors/mids saw 74–83% less coding time on repetitive work; complex tasks flat or negative.	Public-sector pilot	Justifies tier weighting: dense/novel work packs more man-hours per commit.
Verification overhead (reviewing, correcting, cleaning AI output) repeatedly erodes or erases apparent gains.	Recurring across studies	Counted as real time; TDD is treated as billable engineering, not free.

The formula. Per coding day: `man-hours = clamp(commits × 1.6 × tier-weight, 1, 38)`. Days are not fixed 8h — parallel agents push peak days toward 38 man-hours (wall-clock $\leq 18h$), while a light TDD day floors near 1h (~\$25). Actual =

man-hours × tier rate. **Traditional = man-hours × 5 × \$135/hr**: a conventional team, without concurrent agents and absorbing the verification overhead above, reproduces the same tested scope at roughly 5× the labor.

Phase-by-phase build record

Every coding day is a real commit set; every planning day fills a gap. Priced Actual · Billable · Traditional.

General Platform Work

2 coding · 0 planning days

DATE	WORK PRODUCT (FROM COMMITS)	MODEL · TIER	HRS	ACTUAL	BILLABLE 5×	TRADITIONAL
08-11	<code>feat(specs)</code> populate PDP specs via metafields + resilient mapper + Reno extraction <code>feat(specs)</code> harden AI spec extraction, metafield type resilience, and Reno concurrency <code>feat(specs)</code> fix PDP spec rendering for optional fields and wire Shopify metafield read path <code>fix(specs)</code> guard metafield transform against null entries from Storefront API <code>docs(spec)</code> add magnetic track light ingestion design spec	Premium Claude Sonnet 4.5	8.0h	\$400	\$2,000	\$5,400 40.0h
08-12	<code>feat(eiko)</code> migrate product documents S3 to B2 with searchable metadata <code>feat(canolight)</code> add SKU normalization for magnetic track products <code>feat(specs)</code> AI extraction for Eurofase/Votatec + metafield push	Super Premium Claude Opus 4.8	29.1h	\$2,328	\$11,640	\$19,642 145.5h
Subtotal			37.1h	\$2,728	\$13,640	\$25,042

DATE	WORK PRODUCT (FROM COMMITTS)	MODEL · TIER	HRS	ACTUAL	BILLABLE 5×	TRADITIONAL
	pipeline + coverage audit					
	• <code>fix(canolight)</code> remove redundant operations in SKU normalization					
	• <code>feat(canolight)</code> add magnetic track PDF parser					
	• <code>feat(canolight)</code> add magnetic track category derivation					
	• <code>feat(canolight)</code> add --ingest- magnetic-track mode					
	• <code>fix(canolight)</code> export MagneticTrackProduct interface for reusability					
	• <code>fix(canolight)</code> export MagneticTrackProduct interface and fix originalSku to store raw SKU					
	• <code>fix(canolight)</code> improve CRI extraction to avoid matching price numbers					
	• <code>docs(canolight)</code> add magnetic track ingestion plan					
	• <code>docs(quotes)</code> add canolight spec-filler quote (2026-08-12)					
	• <code>feat(pdp)</code> horizontal thumbnail strip with on-image prev/next nav					
Subtotal			37.1h	\$2,728	\$13,640	\$25,042

